

# Skyler's Session

June 29, 2016

```
In [2]: #install.packages(c("caret"))
#install.packages(c("pROC"))
#install.packages(c("gbm"))
#install.packages(c("plyr"))

library(caret)
library(pROC)
library(gbm)
library(plyr)

#####
#####
#####

e <- read.csv("dsss_5.csv")

#The following lines turn the numeric (0/1) label into a FACTOR
DFT<- as.factor(e[,1])
e[,1]<- NULL
levels(DFT) <- list(good = "0", dft = "1")
e<-data.frame(e,DFT)

#####
#####
#####

grid <- expand.grid(interaction.depth = c(1),          #what depth of trees should we grow?
n.trees = seq(5, 150, by=5),          #how many of each depth?
shrinkage = c(.2),          #smaller steps gives better accuracy but requires more trees
n.minobsinnode = 10)          #helps prevent trees from making leaves with few observations

cvCtrl <- trainControl(method = "cv",          #cross validation
number =5,          #how many folds?
summaryFunction = twoClassSummary, #allows us to calculate ROC metrics
classProbs = TRUE, #keep probabilities
verboseIter= TRUE) #watch updating of folds
```

```

gbm_test <- train(DFT ~ .,                                #DFT is the label '.' is short for everything
#Can also try DFT ~ VAR1 + VAR2 to only use 2 predictors
data = e,                                                #Our data
method = "gbm",                                          #use gradient boosted trees, caret has dozens others
distribution = "adaboost", #there are multiple loss functions in GBM, also try "bernoulli"
bag.fraction = .5,                                       #combine boosting and bagging
metric = "ROC",                                         #Choose the best model by maximizing ROC (enabled by twoClassSummary)
trControl = cvCtrl,                                     #perform cross validation according to above values
tuneGrid = grid,                                         #explore depth and tree numbers according to above values
verbose = FALSE,                                         #be quiet
na.action = na.pass)                                     #caret can remove na values, but gbm accepts them.

#####
#####
#####

plot(gbm_test)

plot(gbm_test$finalModel, i.var = 1)

#####
#####
#####

pretty.gbm.tree(gbm_test$finalModel, i.tree = 1)

#####
#####
#####

results <- predict(gbm_test,                             #runs the data set 'e' through the gbm_test model
newdata = e,
type = "prob",      #give probabilities instead of assigning a predicted label
na.action = na.pass) #gbm model can handle na values

b=50                                                         #number of bins for our histograms
g<- results[ e$DFT == "good", ]$good                       #store probability for all the paying customers
d<- results[ e$DFT == "dft" , ]$good                       #store probability for all the defaulting customers

hg <- hist(g, breaks =b, plot = F)                          #make 2 histograms of good and bad, don't plot
dg <- hist(d, breaks =b, plot = F)

top <- max(hg$counts, dg$counts)                            #makes for a prettier graph, stores the highest bin

hist( g,
col=rgb(0,0,1,0.5),                                         #goods are blue-ish with 0.5 transparency

```

```

breaks=b,
xlim = c(0, 1), ylim=c(0,top),      #top of histogram is based on highest of either good or bad
xlab = "Probability of Paid", main = "", ylab = "No. of Customers")


legend("topright", c("Paid","Default"), fill=c( rgb(0,0,1,0.5), rgb(1,0,0,0.5))  )

hist( d,
col=rgb(1,0,0,0.5),                #defaulters are redish with 0.5 transparency
breaks=b,
add=T)                             #add this to the previous histogram instead of making a new

#####
#####
#####


+ Fold1: shrinkage=0.2, interaction.depth=1, n.minobsinnode=10, n.trees=150
- Fold1: shrinkage=0.2, interaction.depth=1, n.minobsinnode=10, n.trees=150
+ Fold2: shrinkage=0.2, interaction.depth=1, n.minobsinnode=10, n.trees=150
- Fold2: shrinkage=0.2, interaction.depth=1, n.minobsinnode=10, n.trees=150
+ Fold3: shrinkage=0.2, interaction.depth=1, n.minobsinnode=10, n.trees=150
- Fold3: shrinkage=0.2, interaction.depth=1, n.minobsinnode=10, n.trees=150
+ Fold4: shrinkage=0.2, interaction.depth=1, n.minobsinnode=10, n.trees=150
- Fold4: shrinkage=0.2, interaction.depth=1, n.minobsinnode=10, n.trees=150
+ Fold5: shrinkage=0.2, interaction.depth=1, n.minobsinnode=10, n.trees=150
- Fold5: shrinkage=0.2, interaction.depth=1, n.minobsinnode=10, n.trees=150
Aggregating results
Selecting tuning parameters
Fitting n.trees = 135, interaction.depth = 1, shrinkage = 0.2, n.minobsinnode = 10 on full training set


```



Out[2]:

	SplitVar	SplitCodePred	LeftNode	RightNode	MissingNode	ErrorReduction	Weight	Prediction
0	4	-0.195	1	2	3	803.7961	5000	0.004228327
1	-1	-0.1003761	-1	-1	-1	0	1794	-0.1003761
2	-1	0.06276243	-1	-1	-1	0	3206	0.06276243
3	-1	0.004228327	-1	-1	-1	0	5000	0.004228327





In [ ]: